

## 0.) Import the Credit Card Fraud Data From CCLE

In [3]:

```
import pandas as pd
from google.colab import drive
import matplotlib.pyplot as plt
import numpy as np
```

In [2]:

```
drive.mount('/content/gdrive/', force_remount = True)
```

Mounted at /content/gdrive/

In [4]:

```
df = pd.read_csv("/content/gdrive/MyDrive/W24ML Code/Data/fraudTest.csv")
```

In [5]:

```
df.head()
```

Out[5]:

	Unnamed: 0	trans_date_trans_time	cc_num	merchant	category	amt	first	last	gender	s
0	0	2020-06-21 12:14:25	2291163933867244	fraud_Kirlin and Sons	personal_care	2.86	Jeff	Elliott	M	Da C
1	1	2020-06-21 12:14:33	3573030041201292	fraud_Sporer-Keebler	personal_care	29.84	Joanne	Williams	F	M l
2	2	2020-06-21 12:14:53	3598215285024754	fraud_Swaniawski, Nitzsche and Welch	health_fitness	41.28	Ashley	Lopez	F	Vale
3	3	2020-06-21 12:15:15	3591919803438423	fraud_Haley Group	misc_pos	60.05	Brian	Williams	M	3 Ki Mil
4	4	2020-06-21 12:15:17	3526826139003047	fraud_Johnston-Casper	travel	3.19	Nathan	Massey	M	F Apt

5 rows × 23 columns

In [6]:

```
df_select = df[["trans_date_trans_time", "category", "amt", "city_pop", "is_fraud"]]

df_select["trans_date_trans_time"] = pd.to_datetime(df_select["trans_date_trans_time"])
df_select["time_var"] = [i.second for i in df_select["trans_date_trans_time"]]

X = pd.get_dummies(df_select, ["category"]).drop(["trans_date_trans_time", "is_fraud"], axis = 1)
y = df["is_fraud"]
```

/var/folders/j8/qj6z29\_s2qj2dwzv274nkt9h0000gp/T/ipykernel\_20135/2282180580.py:3: SettingWithCopyWarning :

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

df\_select["trans\_date\_trans\_time"] = pd.to\_datetime(df\_select["trans\_date\_trans\_time"])

/var/folders/j8/qj6z29\_s2qj2dwzv274nkt9h0000gp/T/ipykernel\_20135/2282180580.py:4: SettingWithCopyWarning :

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

df\_select["time\_var"] = [i.second for i in df\_select["trans\_date\_trans\_time"]]

## 1.) Use scikit learn preprocessing to split the data into 70/30 in out of sample

In [7]:

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

In [8]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = .3)
```

In [9]:

```
X_test, X_holdout, y_test, y_holdout = train_test_split(X_test, y_test, test_size = .5)
```

http://localhost:8888/nbconvert/html/Downloads/Analysis.ipynb?download=false

Page 1 of 6

```
In [10]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
X_holdout = scaler.transform(X_holdout)
```

## 2.) Make three sets of training data (Oversample, Undersample and SMOTE)

```
In [13]: from imblearn.over_sampling import RandomOverSampler
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import SMOTE
```

```
In [14]: ros = RandomOverSampler()
over_X, over_y = ros.fit_resample(X_train, y_train)

rus = RandomUnderSampler()
under_X, under_y = rus.fit_resample(X_train, y_train)

smote = SMOTE()
smote_X, smote_y = smote.fit_resample(X_train, y_train)
```

## 3.) Train three logistic regression models

```
In [15]: from sklearn.linear_model import LogisticRegression
```

```
In [16]: over_log = LogisticRegression().fit(over_X, over_y)

under_log = LogisticRegression().fit(under_X, under_y)

smote_log = LogisticRegression().fit(smote_X, smote_y)
```

## 4.) Test the three models

```
In [17]: over_log.score(X_test, y_test)
```

Out[17]: 0.9051800667002567

```
In [18]: under_log.score(X_test, y_test)
```

Out[18]: 0.9072194630389405

```
In [19]: smote_log.score(X_test, y_test)
```

Out[19]: 0.9042923294704768

```
In [20]: # We see SMOTE performing with higher accuracy but is ACCURACY really the best measure - it is not.
```

## 5.) Which performed best in Out of Sample metrics?

```
In [21]: # Sensitivity here in credit fraud is more important as seen from last class
```

```
In [22]: from sklearn.metrics import confusion_matrix
```

```
In [23]: y_true = y_test
```

```
In [24]: y_pred = over_log.predict(X_test)
cm = confusion_matrix(y_true, y_pred)
cm
```

Out[24]: array([[75226, 7830],
 [ 74, 228]])

```
In [25]: print("Over Sample Sensitivity : ", cm[1,1] /( cm[1,0] + cm[1,1]))
```

Over Sample Sensitivity : 0.7549668874172185

```
In [26]: y_pred = under_log.predict(X_test)
cm = confusion_matrix(y_true, y_pred)
cm
```

Out[26]: array([[75397, 7659],
 [ 75, 227]])

```
In [27]: print("Under Sample Sensitivity : ", cm[1,1] /( cm[1,0] + cm[1,1]))

Under Sample Sensitivity :  0.7516556291390728

In [28]: y_pred = smote_log.predict(X_test)
cm = confusion_matrix(y_true, y_pred)
cm

Out[28]: array([[75152,  7904],
               [   74,   228]])

In [29]: print("SMOTE Sample Sensitivity : ", cm[1,1] /( cm[1,0] + cm[1,1]))

SMOTE Sample Sensitivity :  0.7549668874172185

In [ ]:
```

## 6.) Pick two features and plot the two classes before and after SMOTE.


```
In [ ]: raw_temp = pd.concat([X_train, y_train], axis =1)

In [ ]:

In [ ]: #plt.scatter(raw_temp[raw_temp["is_fraud"] == 0]["amt"], raw_temp[raw_temp["is_fraud"] == 0]["city_pop"],

plt.scatter(raw_temp[raw_temp["is_fraud"] == 1]["amt"], raw_temp[raw_temp["is_fraud"] == 1]["city_pop"])
plt.legend(["Fraud", "Not Fraud"])
plt.xlabel("Amount")
plt.ylabel("Population")

plt.show()
```



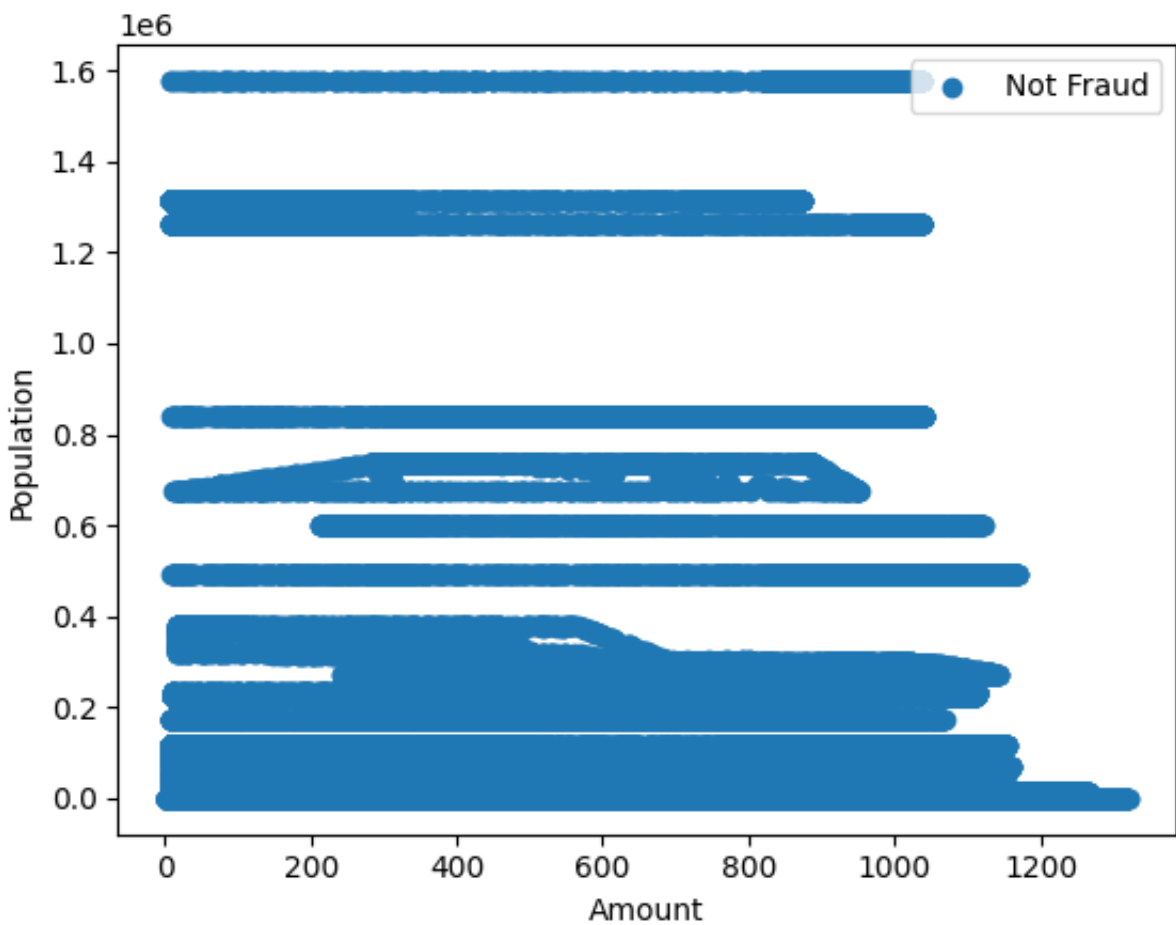
```
In [ ]: raw_temp = pd.concat([smote_X, smote_y], axis =1)

In [ ]: #plt.scatter(raw_temp[raw_temp["is_fraud"] == 0]["amt"], raw_temp[raw_temp["is_fraud"] == 0]["city_pop"],

plt.scatter(raw_temp[raw_temp["is_fraud"] == 1]["amt"], raw_temp[raw_temp["is_fraud"] == 1]["city_pop"])
plt.legend([ "Not Fraud", "Fraud"])
plt.xlabel("Amount")
plt.ylabel("Population")

plt.show()

/usr/local/lib/python3.10/dist-packages/IPython/core/pylabtools.py:151: UserWarning: Creating legend with loc="best" can be slow with large amounts of data.
  fig.canvas.print_figure(bytes_io, **kw)
```



In [ ]:

7.) We want to compare oversampling, Undersampling and SMOTE across our 3 models (LogisticRegression, Logistic Regression Lasso and Decision Trees).

Make a dataframe that has a dual index and 9 Rows.

Calculate: Sensitivity, Specificity, Precision, Recall and F1 score. for out of sample data.

Notice any patterns across performance for this model. Does one totally out perform the others IE. over/under/smote or does a model perform better DT, Lasso, LR?

Choose what you think is the best model and why. test on Holdout

```
In [31]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score
import pandas as pd
```

```
In [33]: resampling_methods = {
    "over": RandomOverSampler(),
    "under":RandomUnderSampler(),
    "smote":SMOTE()
}

model_configs = {
    "LOG":LogisticRegression(),
    "LASSO": LogisticRegression(penalty = "l1", C = 2., solver = "liblinear"),
    "DTREE":DecisionTreeClassifier()
}
```

```
In [35]: trained_models= {}
```

```
In [45]: for resample_key, resampler in resampling_methods.items():
        resample_X, resample_y = resampler.fit_resample(X_train, y_train)

        for model_key, model in model_configs.items():
            combined_key = f"{resample_key}_{model_key}"
            trained_models[combined_key] = model.fit(resample_X, resample_y)
```

```
In [46]: trained_models
```

```
Out[46]: {'over_LOG': LogisticRegression(),
          'over_LASSO': LogisticRegression(C=2.0, penalty='l1', solver='liblinear'),
          'over_DTREE': DecisionTreeClassifier(),
          'under_LOG': LogisticRegression(),
          'under_LASSO': LogisticRegression(C=2.0, penalty='l1', solver='liblinear'),
          'under_DTREE': DecisionTreeClassifier(),
          'smote_LOG': LogisticRegression(),
          'smote_LASSO': LogisticRegression(C=2.0, penalty='l1', solver='liblinear'),
          'smote_DTREE': DecisionTreeClassifier()}
```

```
In [66]: from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score
import pandas as pd

# Initialize an empty list to store results
def calc_perf_metric(y_true, y_pred):
    tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()

    sensitivity = tp / (tp + fn) # Corrected formula
    specificity = tn / (tn + fp) # Corrected formula
    precision = precision_score(y_true, y_pred)
    recall = recall_score(y_true, y_pred) # Recall is the same as sensitivity
    f1 = f1_score(y_true, y_pred)

    # Return a dictionary of the calculated metrics
    return(sensitivity,specificity,precision,recall,f1)
```

```
In [67]: trained_models = {}
results = []
```

```
In [68]: for resample_key, resampler in resampling_methods.items():
        resample_X, resample_y = resampler.fit_resample(X_train, y_train)

        for model_key, model in model_configs.items():
            combined_key = f"{resample_key}_{model_key}"
            m = model.fit(resample_X, resample_y)
            trained_models[combined_key] = m
            y_pred = m.predict(X_test)
            sensitivity,specificity,precision,recall,f1 = calc_perf_metric(y_test, y_pred)
            results.append({"Model": combined_key,
                           'sensitivity': sensitivity,
                           'specificity': specificity,
                           'precision': precision,
                           'recall': recall,
                           'f1_score': f1

                           })

        #results.append(calc_perf_metric(y_true, y_pred))
```

```
In [69]: results_df = pd.DataFrame(results)
```

```
In [70]: results_df
```

Out[70]:

	Model	sensitivity	specificity	precision	recall	f1_score
0	over_LOG	0.751656	0.906750	0.028475	0.751656	0.054871
1	over_LASSO	0.751656	0.906786	0.028485	0.751656	0.054891
2	over_DTREE	0.552980	0.998531	0.577855	0.552980	0.565144
3	under_LOG	0.758278	0.875385	0.021647	0.758278	0.042092
4	under_LASSO	0.761589	0.874795	0.021639	0.761589	0.042082
5	under_DTREE	0.943709	0.949528	0.063659	0.943709	0.119272
6	smote_LOG	0.754967	0.907376	0.028784	0.754967	0.055454
7	smote_LASSO	0.754967	0.907376	0.028784	0.754967	0.055454
8	smote_DTREE	0.725166	0.993896	0.301653	0.725166	0.426070

## Performance Analysis

### 1. Sensitivity (Recall):

- The highest sensitivity is observed with the `under_DTREE` model, indicating it is best at identifying positive cases.
- Both Logistic Regression models ( `over_LOG` , `smote_LOG` ) and the Lasso models show similar sensitivity across the resampling methods, suggesting consistency in identifying true positives regardless of the resampling strategy.

### 2. Specificity:

- The Decision Tree models ( `over_DTREE` , `smote_DTREE` ) exhibit high specificity, particularly with the `over_DTREE` achieving nearly perfect specificity. This indicates strong performance in identifying true negatives.

### 3. Precision:

- The `over_DTREE` model significantly outperforms others in precision, which suggests it has the lowest rate of false positives. This is crucial in scenarios where the cost of a false positive is high.

### 4. F1 Score:

- The `smote_DTREE` model stands out with the highest F1 score, indicating a balanced performance between precision and recall. This balance is important for achieving overall accuracy.

## Pattern Recognition

- **Decision Trees with SMOTE** ( `smote_DTREE` ) shows a compelling balance across all metrics, particularly in achieving a high F1 score, which suggests a balanced trade-off between precision and recall.
- **Oversampling** tends to favor specificity but at the cost of precision, as seen in the `over_DTREE` model.
- **Undersampling** significantly enhances sensitivity, especially for the Decision Tree model ( `under_DTREE` ), making it an excellent choice for applications where missing a positive case (false negative) is critical.
- Logistic Regression models, both standard and with Lasso, display consistency across resampling techniques but do not excel in any particular metric compared to Decision Trees.

## Best Model Selection

Considering the balance between all metrics, **Decision Trees with SMOTE** ( `smote_DTREE` ) appears to be the best model due to its:

- High sensitivity, ensuring most positive cases are identified.
- Excellent specificity, minimizing false positives.
- Strong F1 score, indicating a balanced precision and recall, making it versatile for various applications.

This model strikes a good balance between identifying positive cases without significantly increasing the false positives, making it a robust choice for many scenarios. Testing this model on a holdout dataset would be the next step to validate its performance on unseen data and ensure its generalizability and effectiveness.

In [ ]: